

# Phase 3: Structured Comparison Output

Weekly Exercises — Months 5–6

## Phase Overview

**By the end of this phase, you'll move beyond raw search results to structured comparison data. You'll define Pydantic models that capture what matters when comparing components—props, variants, accessibility features—and use an LLM to extract this information from documentation.**

## What You'll Build

- Pydantic models for component comparisons (props, variants, accessibility)
- An LLM provider interface with modular implementations
- Extraction pipelines that convert documentation into structured data
- Comparison output that normalizes data across design systems

## Key Python Concepts

**Pydantic models and validation, Optional types and defaults, nested models, JSON schema generation, working with LLM APIs, prompt engineering basics, error handling for unreliable outputs.**

## Week 1: Core Comparison Models

### Goal

**Define the Pydantic models that represent component information. Your design system expertise shapes these models—you decide what fields matter for meaningful comparison.**

### Exercise

**Create `src/ds_compare/comparison/models.py` with these Pydantic models:**

- **PropDefinition:** name (str), type (str), required (bool), default (str | None), description (str | None)
- **VariantInfo:** name (str), description (str | None), is\_default (bool)
- **AccessibilityInfo:** aria\_role (str | None), keyboard\_interactions (list[str]), screen\_reader\_notes (str | None)
- **ComponentSummary:** name (str), design\_system (str), description (str | None), props (list[PropDefinition]), variants (list[VariantInfo]), accessibility (AccessibilityInfo | None)

- **ComponentComparison:** component\_name (str), systems (list[ComponentSummary]), generated\_at (datetime)

Use Field() with descriptions. Test by creating instances and verifying validation catches bad data.

## Why This Matters

**These models define what "comparing components" means. Pydantic validates automatically—malformed LLM output gives clear errors, not silent failures. Your domain expertise shapes the schema; a generic developer wouldn't know to include aria\_role or keyboard\_interactions.**

## Completion Checklist

- All five Pydantic models defined with type hints
- Optional used correctly for nullable fields
- Can create valid instances of each model
- Validation rejects invalid data

# Week 2: The LLM Provider Interface

## Goal

**Define an interface for LLM providers, mirroring the embedding pattern. This keeps extraction code decoupled from specific services.**

## Exercise

**Create src/ds\_compare/llm/ with base.py defining:**

- **LLMResponse:** dataclass with content (str), model (str), usage (dict with prompt\_tokens, completion\_tokens)
- **LLMProvider Protocol:** complete(prompt: str, system: str | None, temperature: float) -> LLMResponse
- **complete\_json(prompt: str, schema: type[BaseModel], system: str | None) -> BaseModel:** Returns validated Pydantic model
- **LLMConfig:** dataclass with provider (str), model (str), api\_key (str | None), temperature (float = 0.0)

The complete\_json method is key—it should prompt the LLM for JSON output, parse the response, and validate against the provided Pydantic model.

## Why This Matters

**The complete\_json method encapsulates the entire flow of getting structured data from an LLM: prompting for JSON, parsing the response, validating with Pydantic. This abstraction lets you swap between Claude, GPT-4, or local models without changing extraction code.**

## Completion Checklist

- LLMResponse dataclass defined
- LLMProvider Protocol with both complete methods
- LLMConfig supports multiple providers

- Type hints complete

## Week 3: Anthropic LLM Provider

### Goal

**Implement an LLM provider using the Anthropic API. Claude excels at structured extraction tasks and following detailed instructions.**

### Exercise

**Create `IIm/anthropic.py` with `AnthropicProvider`:**

- Install: pip install anthropic
- Load API key from parameter or ANTHROPIC\_API\_KEY env var
- Implement `complete()` using `client.messages.create()`
- For `complete_json()`: generate JSON schema from Pydantic model using `model.model_json_schema()`
- Include schema in prompt and instruct model to return only valid JSON
- Parse response and validate with the Pydantic model
- Handle JSON parsing errors with clear messages

Test by extracting a simple model (like `PropDefinition`) from a snippet of documentation text.

### Why This Matters

**This is where LLMs become useful for data extraction. The `model_json_schema()` method generates a schema the LLM can follow. Prompt engineering matters here—clear instructions about returning only JSON prevent the model from adding explanatory text that breaks parsing.**

### Completion Checklist

- `AnthropicProvider` implements `LLMProvider`
- API key loaded from env or parameter
- `complete()` returns `LLMResponse`
- `complete_json()` returns validated Pydantic model
- JSON parse errors handled gracefully

## Week 4: OpenAI Provider and Factory

### Goal

**Add an OpenAI provider and create a factory function. This completes the LLM abstraction layer with two provider options.**

### Exercise

**Create `IIm/openai.py` with `OpenAIProvider`:**

- Similar structure to `AnthropicProvider`

- Use response\_format={'type': 'json\_object'} for JSON mode
- Include schema in system prompt for complete\_json()

In llm/\_\_init\_\_.py, implement:

- get\_llm\_provider(config: LLMConfig) -> LLMPromoter factory
- Support 'anthropic' and 'openai' provider values
- Add llm section to your config.yaml

Test both providers with the same extraction task. Compare results quality.

## Why This Matters

**Different LLMs have different strengths. OpenAI's JSON mode guarantees valid JSON (though not schema compliance). Having both options lets users choose based on cost, quality, or preference. The factory pattern makes this transparent to the rest of your code.**

## Completion Checklist

- OpenAIProvider implements LLMPromoter
- Uses JSON mode for reliable parsing
- Factory function returns correct provider
- Config file controls provider selection
- Both providers produce valid Pydantic models

## Week 5: Component Extraction Pipeline

### Goal

**Build the extraction pipeline that converts documentation chunks into ComponentSummary objects. This is where retrieval meets structured output.**

### Exercise

**Create comparison/extraction.py with ComponentExtractor class:**

- \_\_init\_\_ takes LLMPromoter and Searcher
- extract\_component(name: str, design\_system: str) -> ComponentSummary
- Internally: search for relevant chunks, concatenate content, call LLM with extraction prompt

Write a detailed extraction prompt that:

- Explains the task: extract component information from documentation
- Provides the documentation content
- Includes the JSON schema
- Specifies to return ONLY valid JSON, no explanations
- Handles missing information gracefully (use null for unknown fields)

Test by extracting Button from one design system. Verify the output matches the actual documentation.

## Why This Matters

**This is the core of turning unstructured docs into structured data. Prompt engineering is critical—clear instructions reduce extraction errors. Combining search results with LLM extraction is a classic RAG pattern: retrieve context, then generate structured output based on that context.**

## Completion Checklist

- ComponentExtractor class implemented
- Searches for relevant chunks before extraction
- Extraction prompt is clear and detailed
- Returns valid ComponentSummary
- Extracted data matches source documentation

# Week 6: Multi-System Comparison

## Goal

**Build the comparison function that extracts component information from multiple systems and returns a unified ComponentComparison.**

## Exercise

### Add to ComponentExtractor:

- compare\_component(name: str, systems: list[str]) -> ComponentComparison
- Call extract\_component for each system
- Handle cases where a component doesn't exist in a system (skip with warning)
- Assemble results into ComponentComparison with timestamp
- Add progress reporting (extracting from system X...)

Handle component name normalization—the same component might be called 'Dialog' in one system and 'Modal' in another. For now, accept aliases as a parameter: compare\_component(name: str, systems: list[str], aliases: dict[str, str] | None = None).

Test by comparing Button across Spectrum and Carbon. Verify both summaries are populated.

## Why This Matters

**This is the core feature of your tool—comparing components across systems. The aliases parameter acknowledges that naming varies; you'll expand this into smarter normalization later. Graceful handling of missing components prevents the whole comparison from failing when one system lacks a component.**

## Completion Checklist

- compare\_component method implemented
- Extracts from multiple systems
- Handles missing components gracefully
- Alias parameter works for name variations
- Returns valid ComponentComparison with multiple summaries

## Week 7: Validation and Error Handling

### Goal

**Add robust error handling for LLM extraction failures. LLMs are unreliable—sometimes they return malformed JSON or miss required fields. Good error handling makes the difference between a demo and a usable tool.**

### Exercise

**Create comparison/errors.py with custom exceptions:**

- **ExtractionError:** Base exception for extraction failures
- **JSONParseError:** LLM returned invalid JSON
- **ValidationError:** JSON valid but doesn't match schema
- **ComponentNotFoundError:** No relevant chunks found for component

Update ComponentExtractor with retry logic:

- On JSON parse failure, retry up to 3 times with more explicit prompt
- On validation failure, log which fields failed and retry
- After max retries, raise appropriate exception with context
- Log all extraction attempts for debugging

### Why This Matters

**LLMs fail in predictable ways. Sometimes they add explanatory text before JSON. Sometimes they hallucinate fields. Retry logic with clearer prompts often succeeds on the second attempt. Custom exceptions let calling code handle different failure modes appropriately. This is essential for production AI systems.**

### Completion Checklist

- Custom exception classes defined
- Retry logic implemented with max attempts
- Retries use clearer prompts
- Failures logged with context
- Exceptions include useful debugging information

## Week 8: CLI and Integration Test

### Goal

**Add comparison commands to your CLI and run a full integration test. This validates the entire Phase 3 pipeline.**

### Exercise

**Add commands to your CLI:**

- **ds-compare extract <component> --system <name>:** Extract single component, output JSON
- **ds-compare compare <component> --systems <list>:** Compare across systems, output formatted comparison
- **--output <file>:** Optional flag to save JSON to file

For the compare output, create a simple text formatter that shows a side-by-side comparison of props, variants, and accessibility features.

Run integration test: compare Button across Spectrum, Carbon, and Shoelace. Verify:

- All three systems have extracted data
- Props include expected entries (variant, disabled, size, etc.)
- Accessibility info populated where available
- JSON output is valid and parseable

## Why This Matters

**This is the payoff: you can now run a single command to compare how different design systems implement the same component. The JSON output enables programmatic use—feeding into reports, dashboards, or further analysis. The integration test validates that every component from Phase 1–3 works together.**

## Completion Checklist

- extract command works
- compare command works with multiple systems
- JSON output option saves to file
- Text formatter shows readable comparison
- Integration test passes with 3 design systems
- Extracted data is accurate against source docs

## Phase 3 Complete

### By the end of Week 8, you have:

- Pydantic models defining component comparison structure
- Modular LLM provider interface with Anthropic and OpenAI implementations
- Extraction pipeline that converts docs to structured data
- Multi-system comparison capability
- Robust error handling for LLM failures
- CLI commands for extraction and comparison

Before moving to Phase 4, run several comparisons and evaluate extraction quality. Note where the LLM misses information or makes errors—this feedback will inform prompt improvements. Try comparing at least 5 different components across your ingested systems.